

Class 9: Candy Mini-Project

Aadhya Tripathi (PID: A17878439)

Table of contents

Background	1
Data Import	1
Exploratory analysis	3
Overall Candy Rankings	7
Taking a look at pricepercent	14
Exploring the correlation structure	15
Principal Component Analysis	16
Summary	21

Background

In today's mini-project we will analyze candy data with exploratory graphics, ggplot, basic statistics, correlation analysis, and principal component analysis methods we have been learning.

Data Import

The data comes as a CSV file from FiveThirtyEight.

```
candy <- read.csv("candy-data.csv", row.names = 1)
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer
100 Grand	1	0	1			0	0		1
3 Musketeers	1	0	0			0	1		0
One dime	0	0	0			0	0		0
One quarter	0	0	0			0	0		0

Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

There are 85 different candy types.

Q2. How many fruity candy types are in the dataset?

There are 38 fruity candy types.

Q3. What is your favorite candy (other than Twix) in the dataset and what is it's winpercent value?

My favorite candy is “Milky Way” with a winpercent of 73.099556.

Q4. What is the winpercent value for “Kit Kat”?

The winpercent for “Kit Kat” is 76.7686.

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

The winpercent for “Tootsie Roll Snack Bars” is 49.653503.

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

```
library("skimr")
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12

Group variables	None
-----------------	------

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalymond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedrice-wafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

The winpercent variable is on a different scale in comparison to the other columns.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

It is a true/false for whether the candy type is chocolate or not.

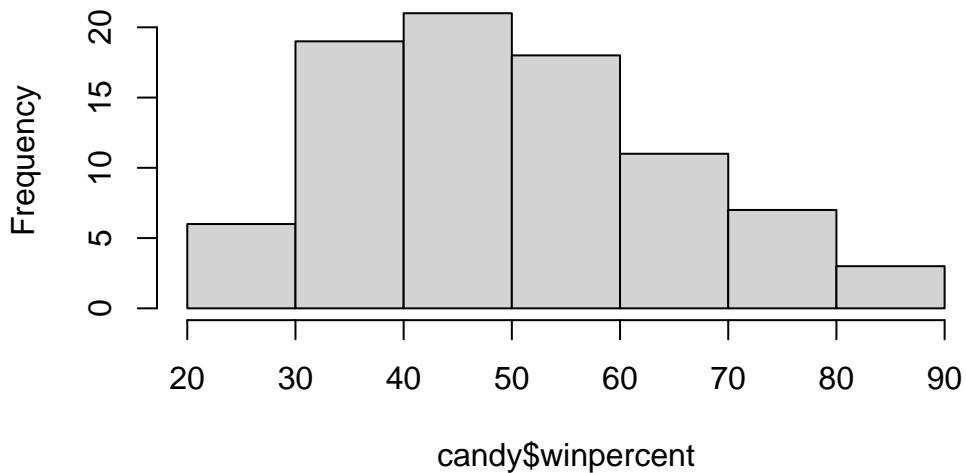
Exploratory analysis

Q8. Plot a histogram of winpercent values using both base R an ggplot2.

Using base R:

```
hist(candy$winpercent)
```

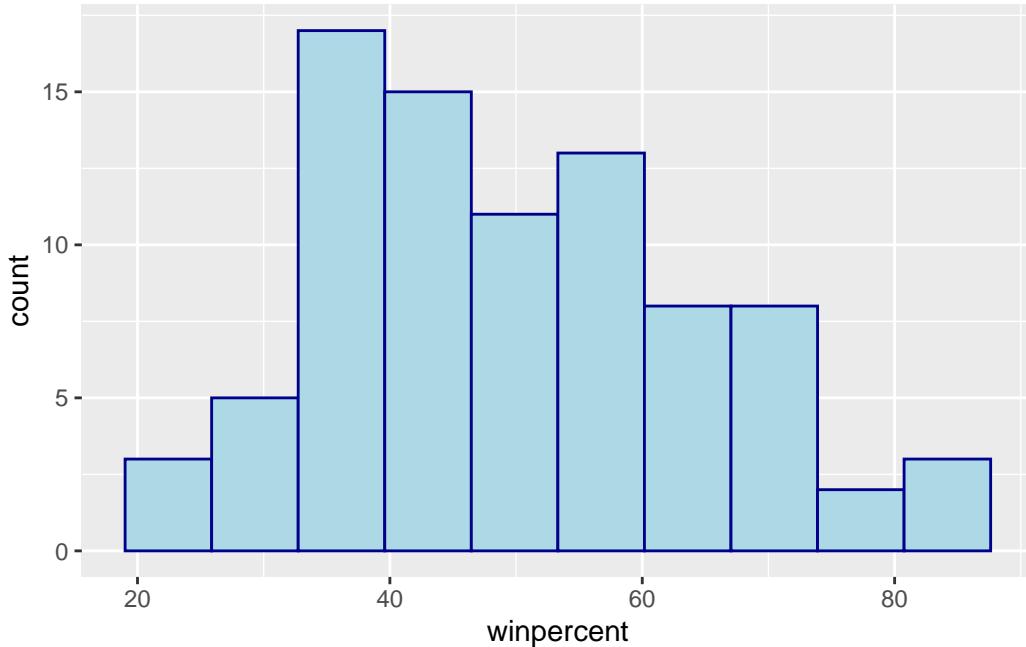
Histogram of candy\$winpercent



Using ggplot:

```
library(ggplot2)

ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins = 10, fill = "lightblue", col = "darkblue")
```



Q9. Is the distribution of winpercent values symmetrical?

No, there is a skew to the right.

Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

[1] 50.31676

```
median(candy$winpercent)
```

[1] 47.82975

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

The mean is slightly above 50%. However, the median is below 50%. Based on the median, the center is determined to be below 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

Steps to solve this problem: 1. Find all chocolate candy in the dataset 2. Find winpercent of these 3. Calculate mean winpercent 4. Repeat 1-3 for fruity candy 5. Compare chocolate mean and fruity mean

```
choc_win <- candy$winpercent[as.logical(candy$chocolate)]
fruity_win <- candy$winpercent[as.logical(candy$fruity)]

print(mean(choc_win))
```

```
[1] 60.92153
```

```
print(mean(fruity_win))
```

```
[1] 44.11974
```

```
mean(choc_win) > mean(fruity_win)
```

```
[1] TRUE
```

Based on their mean values, chocolate candy is ranked higher than fruit candy.

Q12. Is this difference statistically significant?

```
t.test(choc_win, fruity_win)
```

```
Welch Two Sample t-test

data: choc_win and fruity_win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Based on the extremely low p-value, we can reject the null hypothesis and the difference is statistically significant.

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

Use dplyr to rearrange the data by winpercent.

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from 'package:stats':
```

```
filter, lag
```

```
The following objects are masked from 'package:base':
```

```
intersect, setdiff, setequal, union
```

```
candy |>
  arrange(winpercent) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat
Nik L Nip	0	1	0	0	0	0
Boston Baked Beans	0	0	0	1	0	0
Chiclets	0	1	0	0	0	0
Super Bubble	0	1	0	0	0	0
Jawbusters	0	1	0	0	0	0
	crispedrice	wafer	hard	bar	pluribus	sugarpercent
Nik L Nip	0	0	0	1	0.197	0.976
Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

The 5 least liked candy types in this dataset are “Nik L Nip”, “Boston Baked Beans”, “Chiclets”, “Super Bubble”, and “Jawbusters”.

Q14. What are the top 5 all time favorite candy types out of this set?

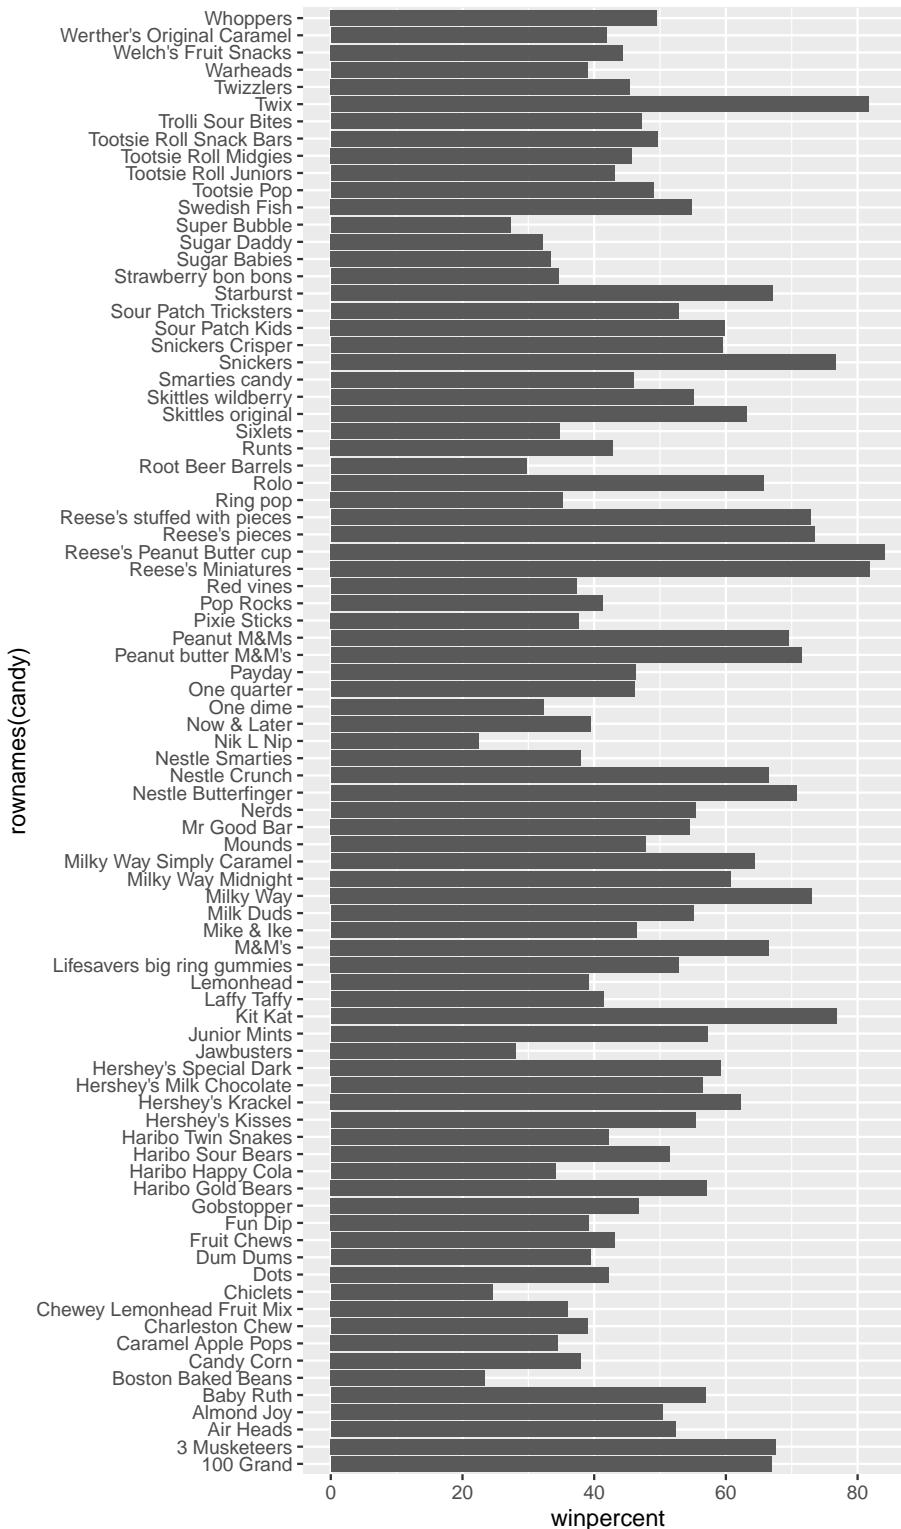
```
candy |>
  arrange(desc(winpercent)) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1
	crispedrice	wafers	hard bar	pluribus	sugar	percent
Reese's Peanut Butter cup	0	0	0	0	0.720	
Reese's Miniatures	0	0	0	0	0.034	
Twix	1	0	1	0	0.546	
Kit Kat	1	0	1	0	0.313	
Snickers	0	0	1	0	0.546	
	price	percent	winpercent			
Reese's Peanut Butter cup	0.651	84.18029				
Reese's Miniatures	0.279	81.86626				
Twix	0.906	81.64291				
Kit Kat	0.511	76.76860				
Snickers	0.651	76.67378				

The 5 all time favorite candy types in this dataset are “Reese’s Peanut Butter cup”, “Reese’s Miniatures”, “Twix”, “Kit Kat”, and “Snickers”.

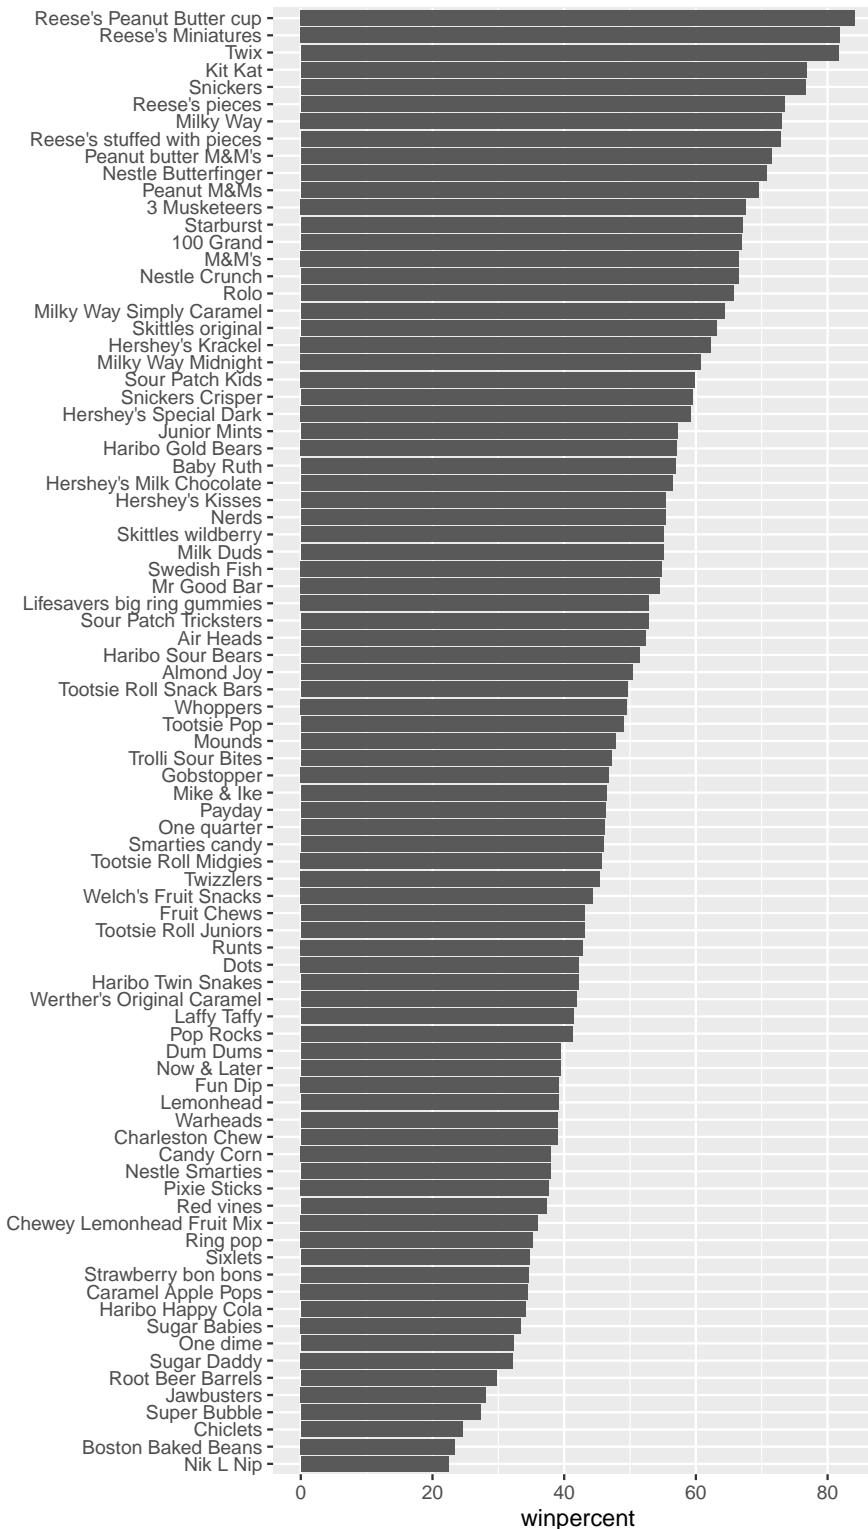
Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent.

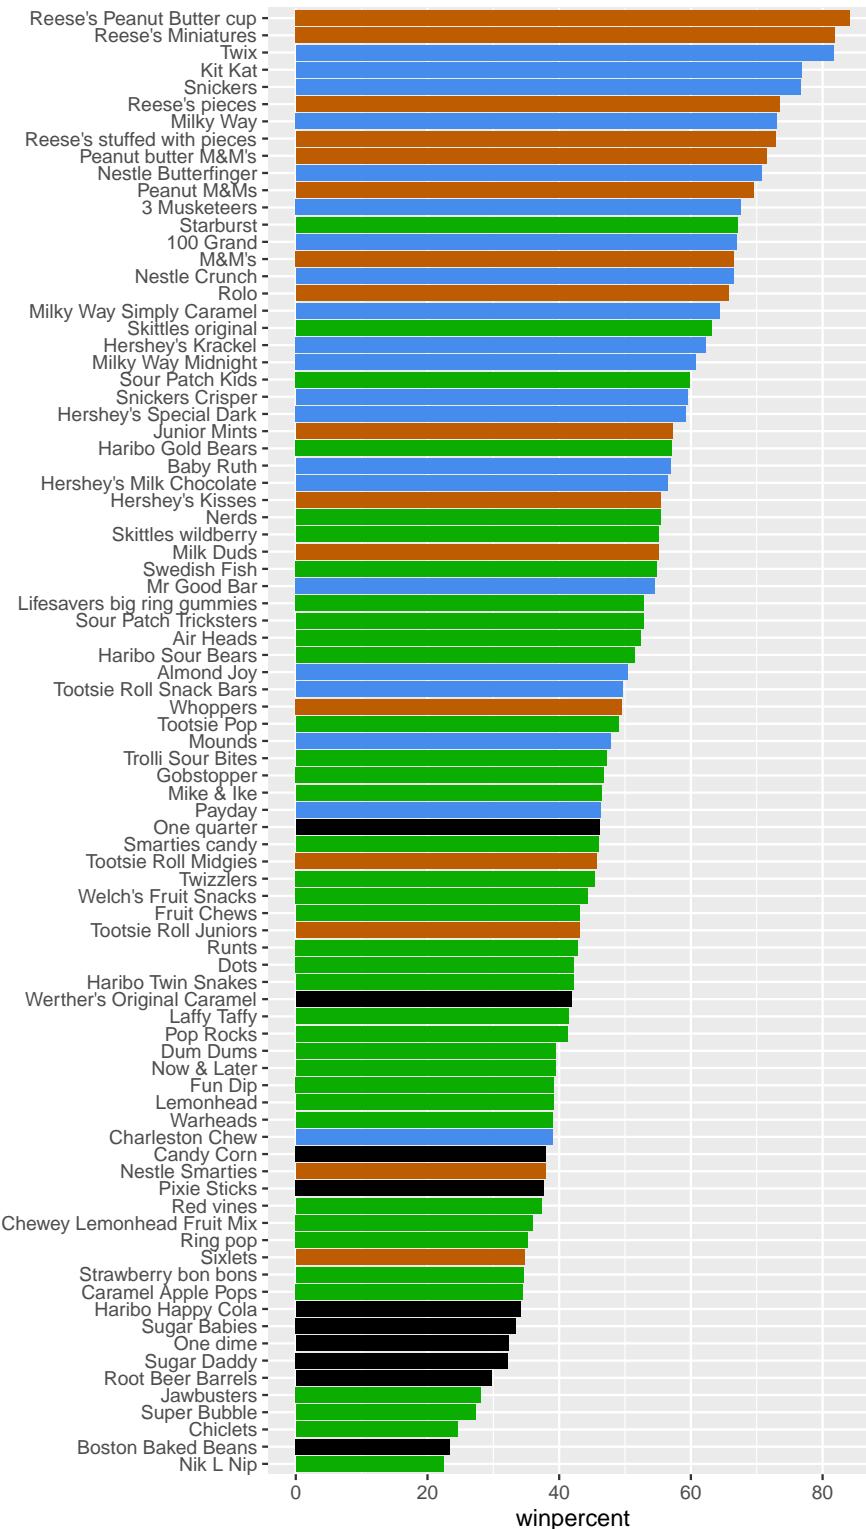
```
ggplot(candy) +  
  aes(winpercent, reorder(rownames(candy),winpercent)) +  
  geom_col() +  
  ylab("")
```



Color the bars based on candy type. Brown is for chocolate, blue for bar candies, and green for fruity candies.

```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "#BD5C00"
my_cols[as.logical(candy$bar)] = "#458CED"
my_cols[as.logical(candy$fruity)] = "#0BAD00"
```

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols) +
  ylab("")
```



Q17. What is the worst ranked chocolate candy?

Sixlets.

Q18. What is the best ranked fruity candy?

Starburst.

Taking a look at pricepercent

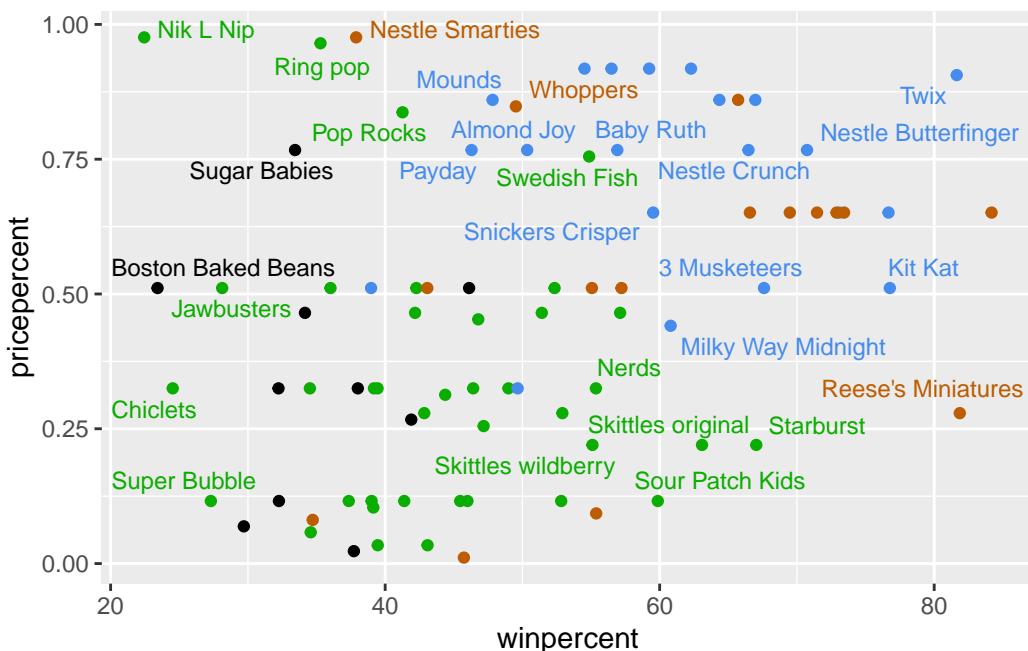
Use `ggrepel` to make overlapping data point labels easier to read.

```
library(ggrepel)
```

Make a plot of winpercent vs pricepercent:

```
ggplot(candy) +  
  aes(winpercent, pricepercent, label=rownames(candy)) +  
  geom_point(col=my_cols) +  
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 7)
```

Warning: ggrepel: 57 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's miniatures have a high winpercent while having a relatively low pricepercent.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
candy |>
  arrange(desc(pricepercent)) |>
  head(5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0	0	0	0
Nestle Smarties	1	0	0	0	0	0
Ring pop	0	1	0	0	0	0
Hershey's Krackel	1	0	0	0	0	0
Hershey's Milk Chocolate	1	0	0	0	0	0
	crisped	rice	wafer	hard	bar	pluribus
Nik L Nip	0	0	0	1	0	0.197
Nestle Smarties	0	0	0	1	0	0.267
Ring pop	0	1	0	0	0	0.732
Hershey's Krackel	1	0	1	0	0	0.430
Hershey's Milk Chocolate	0	0	1	0	0	0.430
	pricepercent	winpercent				
Nik L Nip	0.976	22.44534				
Nestle Smarties	0.976	37.88719				
Ring pop	0.965	35.29076				
Hershey's Krackel	0.918	62.28448				
Hershey's Milk Chocolate	0.918	56.49050				

The top 5 most expensive candy types are “Nik L Nip”, “Nestle Smarties”, “Ring pop”, “Hershey's Krackel”, and “Hershey's Milk Chocolate”. The least popular of these is “Nik L Nip”.

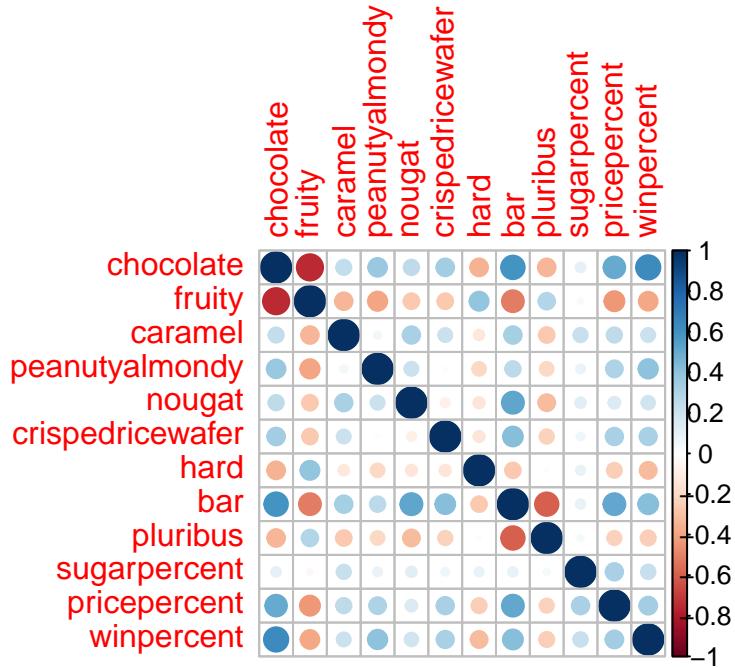
Exploring the correlation structure

Pearson correlation values range from -1 to +1.

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Fruity and chocolate

Q23. Similarly, what two variables are most positively correlated?

Chocolate and winpercent

Principal Component Analysis

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```

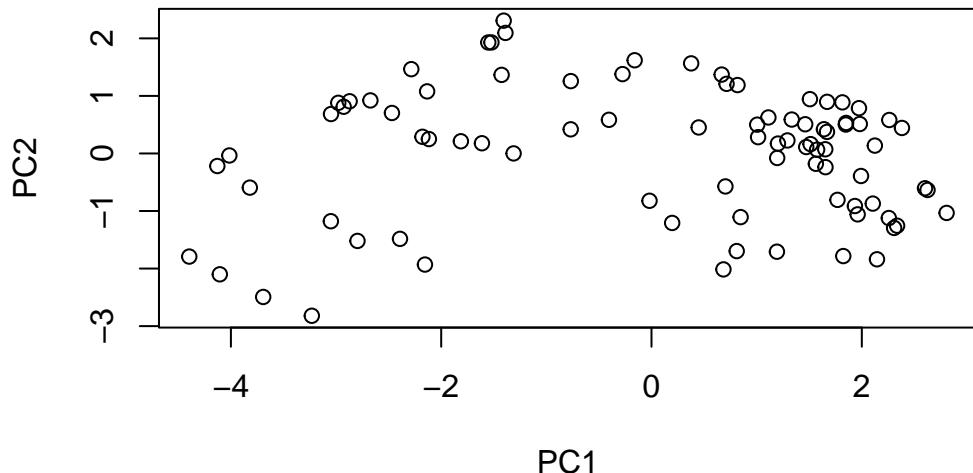
Importance of components:

PC1	PC2	PC3	PC4	PC5	PC6	PC7
-----	-----	-----	-----	-----	-----	-----

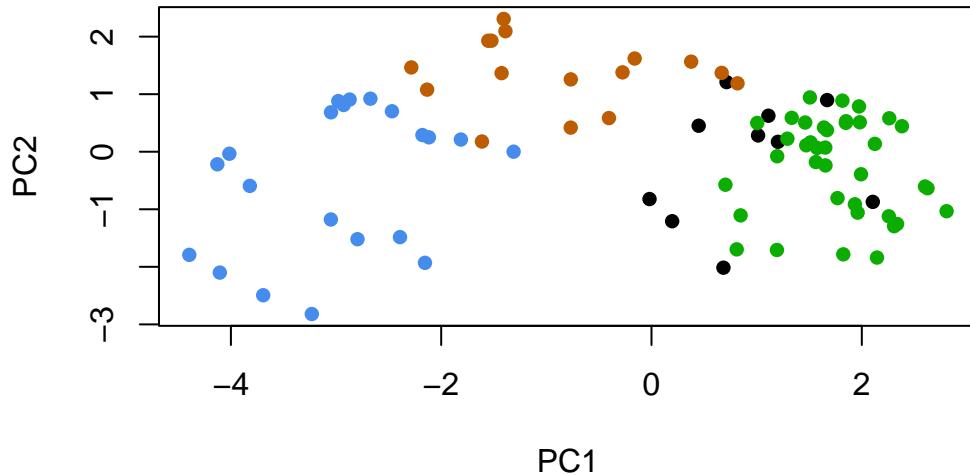
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

The main results figure is the PCA score plot:

```
plot(pca$x[,1:2])
```



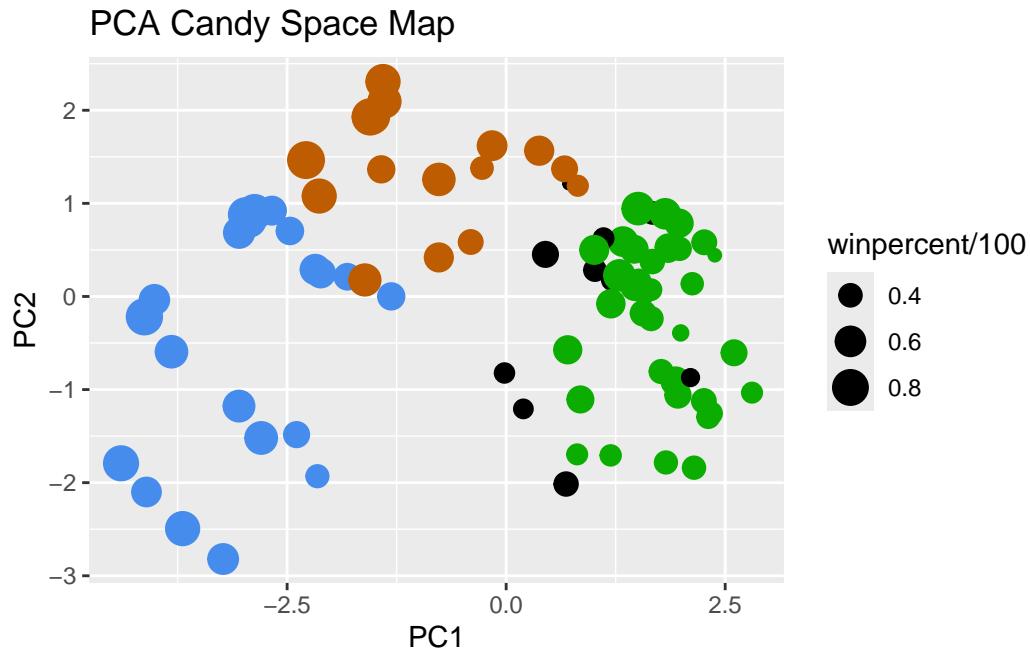
```
plot(pca$x[,1:2], col=my_cols, pch=16)
```



```
# Make a new data-frame with our PCA results and candy data  
my_data <- cbind(candy, pca$x[,1:3])
```

```
p <- ggplot(my_data) +  
  aes(x=PC1, y=PC2,  
      size=winpercent/100,  
      text=rownames(my_data),  
      label=rownames(my_data)) +  
  geom_point(col=my_cols) +  
  labs(title="PCA Candy Space Map")
```

```
p
```



```

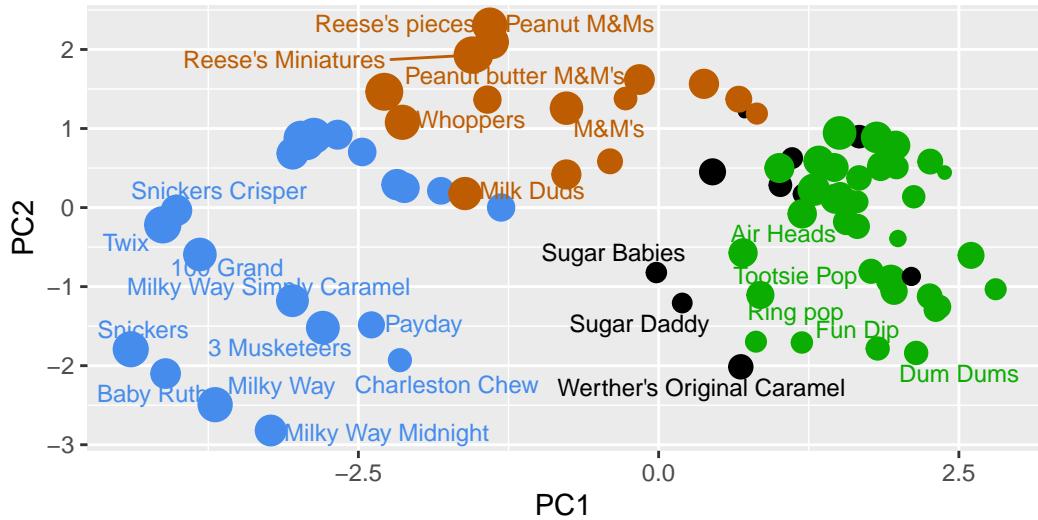
p <- p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (brown), bar (blue), fruity (green), other (black)",
       caption="Data from 538")

```

Warning: ggrepel: 59 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space

Colored by type: chocolate bar (brown), bar (blue), fruity (green), other (black)



Data from 538

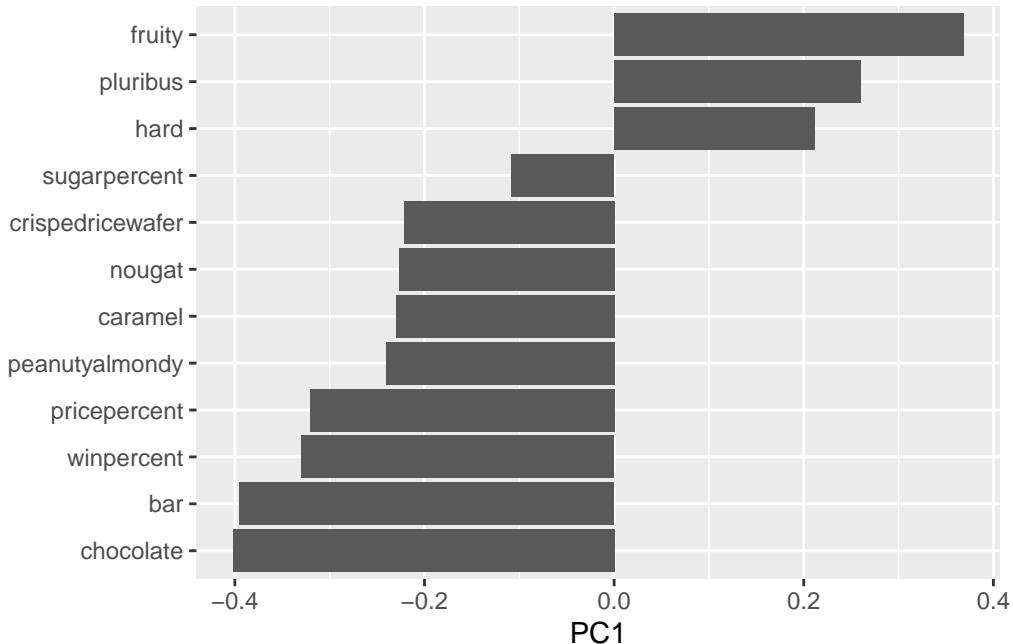
Make an interactive plot with plotly (excluded for PDF render):

```
# library(plotly)
```

```
# ggplotly(p)
```

Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

```
ggplot(pca$rotation) +
  aes(PC1, reorder(rownames(pca$rotation), PC1)) +
  geom_col() +
  ylab("")
```



Fruity, pluribus, and hard are in the positive direction. It makes sense for these characteristics to be correlated as many fruity candies come in multiples and may be hard more often than chocolate. This relationship was previously highlighted in the correlation matrix.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

Chocolate, bar candy types appear to be the “winning” candy. The correlation matrix shows that winpercent has strongest positive correlation with chocolate, and it has a medium positive correlation with bar as well. The PCA plot supports that chocolate and bar candy are more similar to each other compared to fruity candy, as they are closer together on the PC1 axis than chocolate and fruity. The winpercent ordered barplot visualization shows that that brown and blue, representing chocolate and bar candies, tend to be higher compared to the green fruity candies.